

7 things which you should care about before release your code to production

Mateusz Dymiński
Nokia

github.com/mateuszdyminski/7things-java (github.com/mateuszdyminski/7things-java)

Whoami

Mateusz Dymiński:

- Software Developer at Nokia
- 7+ exp with Java
- 3+ exp with Go
- One of the organizer [GoWroc - Golang Wrocław Meetup](https://www.meetup.com/GoWroc) (<https://www.meetup.com/GoWroc>)
- Github: github.com/mateuszdyminski (<http://github.com/mateuszdyminski>)
- Twitter: [@m_dyminski](http://twitter.com/m_dyminski) (http://twitter.com/m_dyminski)
- LinkedIn: linkedin.com/in/mdyminski (<http://linkedin.com/in/mdyminski>)

Agenda

- Versioning
- Profiling
- Health checks
- Logs
- Performance
- Monitoring and alarming
- Release process

1. Versioning

Version Builds

Version Builds

Add information about the build at build time.

- Version
- Last commit
- Build time
- Env vars

Version Builds - how to

Use maven plugin to add information about build

<https://github.com/ktoso/maven-git-commit-id-plugin> (<https://github.com/ktoso/maven-git-commit-id-plugin>)

Version Builds - usage

Get information about the build and last commit from git.

```
<plugins>
  <plugin>
    <groupId>pl.project13.maven</groupId>
    <artifactId>git-commit-id-plugin</artifactId>
    <version>2.2.2</version>
    <executions>
      <execution>
        <id>get-the-git-infos</id>
        <goals>
          <goal>revision</goal>
        </goals>
      </execution>
      <execution>
        <id>validate-the-git-infos</id>
        <goals>
          <goal>validateRevision</goal>
        </goals>
        <phase>package</phase>
      </execution>
    </executions>
  </plugin>
</plugins>
```


Version Builds

Plugin adds git.properties file to jar

```
#Generated by Git-Commit-Id-Plugin
#Fri Jun 30 22:20:44 CEST 2017
git.build.user.email=
git.build.host=md
git.dirty=true
git.remote.origin.url=https\://github.com/mateuszdyminski/7things-java.git
git.closest.tag.name=
git.commit.id.describe-short=ce7b0e9-dirty
git.commit.user.email=dyminski@gmail.com
git.commit.time=29.06.2017 @ 15\:40\:30 CEST
git.commit.message.full=Initial commit
git.build.version=1.0-SNAPSHOT
git.commit.message.short=Initial commit
git.commit.id.abbrev=ce7b0e9
git.branch=master
git.build.user.name=
git.closest.tag.commit.count=
git.commit.id.describe=ce7b0e9-dirty
git.commit.id=ce7b0e9306ad46c011b97d5c51359fc1911a5231
git.tags=
git.build.time=30.06.2017 @ 22\:20\:44 CEST
git.commit.user.name=Mateusz Dyminski
```

Version Builds

We should load it and make available over Health endpoint

```
private GitRepositoryState getGitRepositoryState() {  
    try {  
        Properties properties = new Properties();  
        properties.load(getClass().getClassLoader().getResourceAsStream("git.properties"));  
        return new GitRepositoryState(properties);  
    } catch (IOException e) {  
        throw new RuntimeException(e);  
    }  
}
```

```
@Path("health")  
public class HealthResource {  
  
    @Inject  
    private GitRepositoryState gitRepositoryState;  
  
    @GET  
    @Produces("application/json")  
    public Response getHealth() {  
        return Response.ok().entity(new HealthStatus(gitRepositoryState)).build();  
    }  
}
```

Version Builds - demo

<http://localhost:8090/health> (<http://localhost:8090/health>)

Version artifacts

Version artifacts

Always store build results.

- Nexus, ftp or something
- Github, Bitbucket, Codeplex
- Dockerhub, private Docker registry

Version artifacts - github - how to

```
mvn deploy:deploy-file -Durl=file:///~\m2-repo \  
    -DrepositoryId=some.id \  
    -Dfile=your-artifact-1.0.jar \  
    [-DpomFile=your-pom.xml] \  
    [-DgroupId=org.some.group] \  
    [-DartifactId=your-artifact] \  
    [-Dversion=1.0] \  
    [-Dpackaging=jar] \  
    [-Dclassifier=test] \  
    [-DgeneratePom=true] \  
    [-DgeneratePom.description="My Project Description"] \  
    [-DrepositoryLayout=legacy] \  
    [-DuniqueVersion=false]
```

Version artifacts - docker - how to

```
<plugin>
  <groupId>com.spotify</groupId>
  <artifactId>docker-maven-plugin</artifactId>
  <configuration>
    <imageName>my-image</imageName>
    ...
  </configuration>
  <executions>
    <execution>
      <id>build-image</id>
      <phase>package</phase>
      <goals>
        <goal>build</goal>
      </goals>
    </execution>
    <execution>
      <id>tag-image</id>
      <phase>package</phase>
      <configuration>
        <image>my-image</image>
        <newName>registry.example.com/my-image</newName>
      </configuration>
    </execution>
  </executions>
</plugin>
```

Version artifacts - docker - how to

```
mvn clean package docker:build -DpushImage
```


Version artifacts - github - how to

<https://github.com/aktau/github-release> (https://github.com/aktau/github-release)

```
# upload a file, for example the OSX binary
$ github-release upload \
  --user mateuszdyminski \
  --repo 7things \
  --tag v0.1.0 \
  --name "myapp" \
  --file bin/darwin/myapp
```

Version API

Version API

Two common approach:

- <http://company.com/api/v2/users>

```
GET v2.0/users HTTP/1.1  
Accept: application/json
```

or:

- <http://company.com/api/users>

```
GET /users HTTP/1.1  
Accept: application/vnd.usersapp-v2+json
```

or:

- <http://company.com/api/users>

```
GET /users HTTP/1.1  
x-app-version: 2.0
```

Version API - path-based

- <http://company.com/api/v2/users>

```
GET v2/users HTTP/1.1 Accept: application/json
```

Code:

```
@Path("/api")
public class UsersResource {

    @GET
    @Path("/v{version}/users")
    @Produces("application/json")
    public Response getUsers(@PathParam("version") String version) {
        if (version.equalsIgnoreCase("1")) {
            return Response.ok()
                .entity(Arrays.asList(new User("Jan", "Kowalski")))
                .build();
        } else {
            return Response.ok()
                .entity(Arrays.asList(new User("Jan", "Kowalski", "987654321")))
                .build();
        }
    }
}
```

Version API - how to

Test:

<http://localhost:8180/api/v1/users> (http://localhost:8180/api/v1/users)

<http://localhost:8180/api/v2/users> (http://localhost:8180/api/v2/users)

Version API - headers - how to

- <http://company.com/api/users>

```
GET /users HTTP/1.1 Accept: application/vnd.usersapp-v2+json
```

Code:

```
@Resource
@Path("users")
public class UsersResource {

    @GET
    @Produces("application/vnd.usersapp-v1+json")
    public Response getUsersV1() {
        User user = new User("Jan", "Kowalski");
        return Response.ok().entity(Arrays.asList(user)).build();
    }

    @GET
    @Produces("application/vnd.usersapp-v2+json")
    public Response getUsersV2() {
        User user = new User("Jan", "Kowalski", "987654321");
        return Response.ok().entity(Arrays.asList(user)).build();
    }
}
```

Version API - headers - how to

Test:

```
curl http://localhost:8180/users
```

```
curl -H "Accept: application/vnd.usersapp-v1+json" http://localhost:8180/users
```

```
curl -H "Accept: application/vnd.usersapp-v2+json" http://localhost:8180/users
```

Version Static files

Version Static files

Reduce the risk that client gets cached-old files.

- /statics/\$REVISION/app.js
- /statics/app.js?v=\$REVISION
- /statics/app.\$REVISION.js

Version DB schema

Version DB schema

- Automatic - one command to run all migrations
- Reversible - rollbacks

Options:

- Flyway
- Liquibase
- Custom solution

Version DB - Liquibase example

changelog.sql

```
--liquibase formatted sql

--changeset mdyminski:1 dbms:mssql
--comment create users table
create table users (
    id int primary key,
    name varchar(255)
);
--rollback drop table users;

--changeset mdyminski:2 dbms:mssql
--comment alter table users - add new field
alter table users add phone nvarchar (255) NULL;
--rollback ALTER TABLE users DROP COLUMN phone;
```

it also might be xml, yaml, json

Version DB - how to run

```
# /bin/bash

liquibase --changeLogFile=changelog.sql
--driver=com.microsoft.sqlserver.jdbc.SQLServerDriver
--classpath=jdbc/sqljdbc4-4.0.jar
--url="jdbc:sqlserver://somedb.database.windows.net:1433;database=some-db"
--username=some-user
--password=some-password
update
```

it also might be migrate, validate, diff, sql output

2. Profiling

Profiling

Add a way to profile your application in any time.

Sometimes you have to switch something on to be able to profile your application during the normal production shift.

Profiling

Options:

- Your Kit
- Java Mission Control
- JProfiler
- VisualVM
- NetBeans profiler
- Jprobe

Profiling

To enable remote access to our Java app add flags:

```
-Dcom.sun.management.jmxremote  
-Dcom.sun.management.jmxremote.port=8011  
-Dcom.sun.management.jmxremote.ssl=false  
-Dcom.sun.management.jmxremote.authenticate=false
```

Don't forget about the firewall/security policy

Profiling

Don't forget about the


```
-XX:+HeapDumpOnOutOfMemoryError -XX:HeapDumpPath=/some/place/dumps
```

3. Health checks

Health checks


Do I need health check ?

Health checks



Health checks

Do I need health checks in following architecture?



Healthz

- It leverage health endpoint pattern.
- Checks connection to the DB.
- Prints much more information about the service health like:

- Build info
- Uptime
- Hostname
- Db connection status

Inspiration:

[GitHub - app-healthz](https://github.com/kelseyhightower/app-healthz) (https://github.com/kelseyhightower/app-healthz)

[Kelsey Hightower - healthz](https://vimeo.com/173610242) (https://vimeo.com/173610242)

Health checks - how to

We could reuse pattern from Versioning Builds:

```
private GitRepositoryState getGitRepositoryState() {  
    try {  
        Properties properties = new Properties();  
        properties.load(getClass().getClassLoader().getResourceAsStream("git.properties"));  
        return new GitRepositoryState(properties);  
    } catch (IOException e) {  
        throw new RuntimeException(e);  
    }  
}
```


Health checks - how to

```
@Path("api/health")
public class ExtendedHealthResource {

    @Inject
    private GitRepositoryState gitRepositoryState;

    @Inject
    private MongoDBDatabase mongoDatabase;

    @Inject
    private AppStatus appStatus;

    @GET
    @Produces("application/json")
    public Response getHealth() {
        Document serverStatus = mongoDatabase.runCommand(new Document("serverStatus", 1));
        return Response.ok()
            .entity(new ExtendedHealthStatus(gitRepositoryState, serverStatus, appStatus))
            .build();
    }
}
```

Health checks - how to

```
public class AppStatus {  
  
    private String hostname;  
    private LocalDateTime StartedAt;  
  
    public AppStatus(String hostname, LocalDateTime startedAt) {  
        this.hostname = hostname;  
        StartedAt = startedAt;  
    }  
  
    public String getHostname() {  
        return hostname;  
    }  
  
    @JsonProperty  
    public String getStartedAt() {  
        return StartedAt.toString();  
    }  
  
    @JsonProperty  
    public String uptime() {  
        return Duration.between(StartedAt , LocalDateTime.now()).toString();  
    }  
}
```

Health checks - how to

```
private AppStatus getAppStatus() {
    String hostname = "Unknown";

    try
    {
        InetAddress addr;
        addr = InetAddress.getLocalHost();
        hostname = addr.getHostName();
    }
    catch (UnknownHostException ex)
    {
        throw new RuntimeException(ex);
    }

    return new AppStatus(hostname, LocalDateTime.now());
}
```

Health checks - Demo

Open

<http://localhost:8090/api/health> (<http://localhost:8090/api/health>)

4. Logs

Logs

- Use structured logger
- Log context of invocation
- Use log aggregators
- Log with error level should force user (administrator, or direct user) intervention!

Logs - simple example

```
@Path("calculator")
public class Calculator {

    private static Logger logger = LogManager.getLogger(Calculator.class);

    @GET
    @Path("add")
    @Produces("text/plain")
    public String add(@QueryParam("val1") Integer val1,
                     @QueryParam("val2") Integer val2) {
        if (val1 != null) {
            logger.info("Got val1: {}", val1);
        }

        if (val2 != null) {
            logger.info("Got val2: {}", val2);
        }

        int result = val1 + val2;
        logger.info("Result: {}", result);

        return format("Result of %d + %d = %d", val1, val2, result);
    }
}
```

Logs - demo

Run

```
curl 'http://localhost:8070/calculator/add?val1=3&val2=7'
```

Got result

Result of $3 + 7 = 10$

Logs

```
[INFO] 22:31:54.649 LF.filter - HTTP REQUEST: GET /calculator/add Headers: {host=[localhost:8070], user-  
[INFO] 22:31:54.653 C.add - Got val1: 3  
[INFO] 22:31:54.654 C.add - Got val2: 7  
[INFO] 22:31:54.655 C.add - Result: 10  
[INFO] 22:31:54.656 LF.filter - HTTP RESPONSE: GET /calculator/add Status: 200
```


Logs - demo v2

Run

```
hey -n 100 -c 10 -H "Content-type: text/plain" "http://localhost:8070/calculator/add?val1=3&val2=7"
```

Logs

```
[INFO] 22:36:24.458 LF.filter - HTTP RESPONSE: GET /calculator/add Status: 200
[INFO] 22:36:24.458 C.add - Got val2: 7
[INFO] 22:36:24.459 C.add - Result: 10
[INFO] 22:36:24.459 LF.filter - HTTP RESPONSE: GET /calculator/add Status: 200
[INFO] 22:36:24.461 LF.filter - HTTP REQUEST: GET /calculator/add Headers: {host=[localhost:8070], user
[INFO] 22:36:24.461 LF.filter - HTTP REQUEST: GET /calculator/add Headers: {host=[localhost:8070], user
[INFO] 22:36:24.462 C.add - Got val1: 3
[INFO] 22:36:24.462 C.add - Got val2: 7
[INFO] 22:36:24.462 C.add - Got val1: 3
[INFO] 22:36:24.462 C.add - Result: 10
[INFO] 22:36:24.462 C.add - Got val2: 7
[INFO] 22:36:24.462 LF.filter - HTTP RESPONSE: GET /calculator/add Status: 200
[INFO] 22:36:24.463 C.add - Result: 10
[INFO] 22:36:24.463 LF.filter - HTTP RESPONSE: GET /calculator/add Status: 200
```

Logs - structured logger + context

```
@PreMatching
public class LoggingFilterWithCtx implements ContainerRequestFilter, ContainerResponseFilter {
    private static Logger logger = LogManager.getLogger(LoggingFilter.class);

    @Inject
    private Provider<Request> request;

    @Override
    public void filter(ContainerRequestContext requestContext) throws IOException {
        ThreadContext.put("requestId", UUID.randomUUID().toString());
        ThreadContext.put("requestIp", request.get().getRemoteAddr());
        logger.info("HTTP REQUEST: {} /{} Headers: {}",
            requestContext.getMethod(), requestContext.getUriInfo().getPath(),
            requestContext.getHeaders());
    }

    @Override
    public void filter(ContainerRequestContext requestContext, ContainerResponseContext responseContext)
        logger.info("HTTP RESPONSE: {} /{} Status: {}",
            requestContext.getMethod(), requestContext.getUriInfo().getPath(),
            responseContext.getStatus());
        ThreadContext.remove("requestId");
        ThreadContext.remove("requestIp");
    }
}
```

Logs - demo v3

Run

```
hey -n 100 -c 10 -H "Content-type: text/plain" "http://localhost:8080/calculator/add?val1=3&val2=7"
```

Logs

```
[I] 22:39:16.19 98c3d2f1-815e216e0c63 [grizzly-11] 127.0.0.1 - HTTP RESPONSE: GET /calculator/add Statu
[I] 22:39:16.19 aab9e772-f9abdf62a129 [grizzly-9] 127.0.0.1 - Got val1: 3
[I] 22:39:16.19 aab9e772-f9abdf62a129 [grizzly-9] 127.0.0.1 - Got val2: 7
[I] 22:39:16.20 aab9e772-f9abdf62a129 [grizzly-9] 127.0.0.1 - Result: 10
[I] 22:39:16.19 a589072c-fa04f4fa2a3b [grizzly-12] 127.0.0.1 - Got val2: 7
[I] 22:39:16.20 a589072c-fa04f4fa2a3b [grizzly-12] 127.0.0.1 - Result: 10
[I] 22:39:16.20 aab9e772-f9abdf62a129 [grizzly-9] 127.0.0.1 - HTTP RESPONSE: GET /calculator/add Status
[I] 22:39:16.20 9568307d-9ace8480f489 [grizzly-14] 127.0.0.1 - HTTP REQUEST: GET /calculator/add Header
[I] 22:39:16.21 9568307d-9ace8480f489 [grizzly-14] 127.0.0.1 - Got val1: 3
[I] 22:39:16.22 9568307d-9ace8480f489 [grizzly-14] 127.0.0.1 - Got val2: 7
[I] 22:39:16.22 9568307d-9ace8480f489 [grizzly-14] 127.0.0.1 - Result: 10
[I] 22:39:16.22 9568307d-9ace8480f489 [grizzly-14] 127.0.0.1 - HTTP RESPONSE: GET /calculator/add Statu
[I] 22:39:16.21 29a0e614-d536ec677462 [grizzly-13] 127.0.0.1 - HTTP REQUEST: GET /calculator/add Header
[I] 22:39:16.22 a589072c-fa04f4fa2a3b [grizzly-12] 127.0.0.1 - HTTP RESPONSE: GET /calculator/add Statu
[I] 22:39:16.23 29a0e614-d536ec677462 [grizzly-13] 127.0.0.1 - Got val1: 3
[I] 22:39:16.24 29a0e614-d536ec677462 [grizzly-13] 127.0.0.1 - Got val2: 7
[I] 22:39:16.24 29a0e614-d536ec677462 [grizzly-13] 127.0.0.1 - Result: 10
```

Log things in the frontend/mobile app

- Add extra endpoint in the backend for logs from frontend - tricky and risky
- Use Sentry or other similar tool
- Sometimes backend works as expected but end user is suffering

Use log aggregators

In case where you would like to have more than 1 node or you don't know the flags for grep :
)

Improve security - hacker can clean up local files but can't remove logs which are already sent to the log aggregators.

Options:

- elasticsearch
- splunk - if you're rich enough
- greylog
- loggly and many more

Use log aggregators - ELK -how to

Configure log4j to send logs to LogStash - old way

```
log4j.rootLogger=debug,tcp  
  
log4j.appender.tcp=org.apache.log4j.net.SocketAppender  
log4j.appender.tcp.Port=5000  
log4j.appender.tcp.RemoteHost=localhost  
log4j.appender.tcp.ReconnectionDelay=10000  
log4j.appender.tcp.Application=playground
```

Use log aggregators - ELK -how to

Use <http://www.fluentd.org/> (<http://www.fluentd.org/>)

```
<source>
  @type tail
  format apache
  tag apache.access
  path /var/log/httpd/access_log
</source>

<match var.log.containers.log**>
  type elasticsearch
  log_level info
  host elasticsearch-logging
  port 9200
  logstash_format true
  num_threads 8
</match>
```

Use log aggregators - ELK -how to

Use FileBeats <https://www.elastic.co/products/beats/filebeat> (<https://www.elastic.co/products/beats/filebeat>)

```
# filebeat.yml
filebeat:
  prospectors:
    -
      paths:
        - /var/log/your-app/app.*.log
      input_type: log
  output:
    logstash:
      hosts: ["your-logstash-host:5000"]
```


5. Performance

Performance

- Run performance tests at least once and save the results.
- Run stability tests at least once.
- Don't forget to run them on environment similar to production.
- This isn't as much of time consuming as many think.

Options:

- <http://gatling.io/> (<http://gatling.io/>)
- <http://jmeter.apache.org/> (<http://jmeter.apache.org/>)
- <https://github.com/tsenart/vegeta> (<https://github.com/tsenart/vegeta>)
- <https://github.com/rakyll/boom> (<https://github.com/rakyll/boom>) / <https://github.com/rakyll/hey>
(<https://github.com/rakyll/hey>)

Performance - report

```
hey -c 10 -n 1000 http://localhost:8090/api/health
```

Result:

Summary:

```
Total:      4.1074 secs
Slowest:    0.2727 secs
Fastest:    0.0353 secs
Average:    0.0389 secs
Requests/sec: 243.4658
```

Status code distribution:

[200] 1000 responses

Response time histogram:

[illegible]

6. Monitoring and alarming

Monitoring and alarming

- Monitoring should never require a human to interpret any part of the alerting domain
- Three valid kinds of monitoring output

- Alerts: human needs to take action immediately
- Tickets: human needs to take action eventually
- Logging: no action needed

- Eliminating toil: Carla Geisser: "If a human operator needs to touch your system during normal operations, you have a bug. The definition of normal changes as your systems grow."

- If you have to ssh to your server to do some work - something is wrong
- Automate backups
- Automate logs rotating - send old ones to S3/equivalent

Monitoring

- Measure everything
- Set alerts based on the metrics
- Analytics can use your metrics!

Monitoring - key things to measure

- Java GC metrics
- Http code rates
- Http response time percentiles - 50, 95, 99
- Error logs
- Uptime
- Business metrics
- IPC(instructions per cycle) over CPU - [CPU utilization is wrong](http://www.brendangregg.com/blog/2017-05-09/cpu-utilization-is-wrong.html) (<http://www.brendangregg.com/blog/2017-05-09/cpu-utilization-is-wrong.html>)

Monitoring - Librato - how to

```
MetricRegistry registry = environment.metrics();  
Librato.reporter(registry, "<Librato Email>", "<Librato API Token>")  
    .setSource("<Source Identifier (usually hostname)>")  
    .start(10, TimeUnit.SECONDS);
```

```
Librato.metric(registry, "logins").tag("uid", uid).meter().mark()  
Librato.metric(registry, "kafka-read-latencies").tag("broker", broker).histogram().update(latency)  
Librato.metric(registry, "temperature").source("celcius").tag("type", "celcius").gauge(() -> 42)  
Librato.metric(registry, "jobs-processed").source("ebs").meter().mark()  
Librato.metric(registry, "just-these-tags").tag("foo", "bar").doNotInheritTags().timer.update(time)
```


7. Release process

Release process - release notes

Create release notes

- JIRA - [How to get release notes from JIRA](https://confluence.atlassian.com/adminjiraserver071/creating-release-notes-802592502.html) (https://confluence.atlassian.com/adminjiraserver071/creating-release-notes-802592502.html)
- GIT - `git log v1.0..v1.2`

Create release plan

- When release the application (time, day)
- What if something goes wrong? Revert?

Release process - Add release automation

- Ansible/Puppet/Chef/Salt*
- Docker/Kubernetes/Swarm
- Sometimes bash is good enough :)
- Use process daemons - Does my app be up & running when node will be rebooted?

Use: systemd/supervisor/runit

Release process - no downtime(or minimize it)

- Graceful shutdown
- Canary deployments
- Rollout deployments - kubernetes

Thank you

Mateusz Dymiński

Nokia

github.com/mateuszdyminski/7things-java (github.com/mateuszdyminski/7things-java)

[@m_dyminski](https://twitter.com/m_dyminski) ([http://twitter.com/m_dyminski](https://twitter.com/m_dyminski))

